

Implementation Plan — Phased by Capability

The deterministic core ships first and proves itself, then the LLM enablement layer, then the decision layer and full team adoption. Implementation isn't done when the code runs — it's done when the team runs on it.

Phase 0 — Build COMPLETE

Blocks 1–7 per `DESIGN.md` §5.

Exit: live [Trigger.dev run link](#), `npm run check` **green**.

Phase 1 — Core pipeline live: enrichment · ICP scoring · routing ZERO-LLM PATH

Week 1 — no model, no vendor key required to deliver value.

- Deploy the core tasks to Trigger.dev prod; env vars set (never committed).
- Enrichment cascade live (vendors env-gated, probe-first; DNS/RDAP + symbolic floor).
- Create **#lead-pipeline-ops** in Slack: run summaries + threshold alerts (unrouted lead, tier skew, rep imbalance, enrichment trust).
- Run against the real book; verify invariants live: zero unrouted, breakdown on every score, `dq_reason` on every DQ, idempotent re-run (ownership ledger stable).
- Ops owner reviews the routed book + run stats; routing disputes → config edits.

Exit: routed book delivered with reasons on every decision; re-run produces the identical book; ops sees alerts in Slack.

Phase 2 — Enablement live: battle cards FIRST LLM CAPABILITY

Week 2 — gated, eval'd, and rep-facing.

- `ANTHROPIC_API_KEY` + Langfuse keys into prod env; groundedness eval active — a card citing anything outside `docs/grounding/` fails before a rep sees it.
- Per-rep Slack delivery: P1 hot-lead alert carries score, breakdown, reason, and battle-card link — **the alert is where the SLA clock starts**.
- **Slack verification — don't assume they see it, prove it:** every rep joins their channel, notifications checked (not muted, mobile push on), and each rep  -reacts to a **test P1 alert**. No ack, no go-live for that book.

- **Training session (45 min, recorded):** how a routed lead arrives · reading the score `breakdown` and `dq_reason` · reading a battle card in 60 seconds · SLAs (P1 same-day, P2 48h, P3 nurture) · **accept/decline with a reason** — the decline reason is structured data that retunes the rubric, so "junk" is not a reason.
- Rep cheat-sheet (one page): tiers, priorities, where outputs live, dispute path.
- **Shadow week:** alerts + cards flow, no quota consequences; collect confusion points, fix config/copy, not habits.

Exit: every rep has ack'd a test alert; cards passing eval; first real P1 touched within SLA.

Phase 3 — Decision layer & adoption

Weeks 3–4.

- Decision brief live: deltas vs. prior run, enrichment trust score, projections, recommendations; LLM narration on (numbers computed by code only).
- **Adoption metrics:** alert-ack rate · time-to-first-touch vs. SLA · ledger fill rate (a decline without a reason counts as unfilled).
- Weekly 15-min review (first month) runs on the brief, not on opinions; changes ship as shadow-diffed config edits per `OPERATIONS.md`.
- Then hand off to the standing `OPERATIONS.md` cadence (weekly / monthly / quarterly).

Instrumentation per phase — what's tracked, and how we'd know it's working

Same metrics as the standing plan (`OPERATIONS.md` §1), staged by when they first matter:

Phase	What's tracked	"It's actually working" when
1 — Core	Per-stage success/retries/duration (Trigger.dev) · tier distribution + rep load vs. alert thresholds · enrichment trust score (% live-signal vs. fallback)	30/30 routed with a rep + score + human-readable reason, zero unrouted, re-run produces the identical book, all thresholds quiet
2 — Cards	Langfuse cost/tokens/latency per card · groundedness eval pass-rate · rep alert-ack rate	100% of shipped cards pass eval, every rep has ack'd, first real P1 touched within SLA
3 — Adoption	Time-to-first-touch vs. SLA by priority · ledger fill rate · reply/meeting rate by tier	≥90% P1 within SLA, 100% of declines carry reasons, and hot outperforms warm

— if it doesn't, the rubric is decoration
and Phase 3 feeds the recalibration loop

The tier-conversion test is the plan's real success metric: everything before it proves the machine runs; that one proves the scoring means something.

Roles

Role	Owns
Growth engineer	Pipeline, config changes, shadow-diff promotion, vendor health
Ops owner	Slack alerting, SLA tracking, adoption metrics, weekly review
Reps	Their book: touch within SLA, accept/decline with reasons
Escalation	Routing dispute → ops owner; unresolved >48h → growth engineer edits config

Definition of "implemented"

1. Core: routed book delivered, zero unrouted, idempotent re-runs verified in prod.
2. Every rep has  -ack'd a test alert in their own channel (seen ≠ assumed).
3. ≥90% of P1 alerts touched within SLA from week 3 on.
4. 100% of declines carry a reason in the ledger.
5. Weekly review runs on the decision brief, not on opinions.

Socure Growth Engineering Exercise · source of truth: [docs/IMPLEMENTATION.md](#)